

# HA2: collection indexing using Elasticsearch (part 1)

- Download elasticsearch  
<https://www.elastic.co/downloads/elasticsearch>
- Install Python Elasticsearch client  
<https://elasticsearch-py.readthedocs.io/>
- Read documentation  
<https://www.elastic.co/guide/en/elasticsearch/reference/8.6/>

# Steps to build a search engine

- Define configuration
- Create empty index
- Index collection
- Perform search
- (Evaluate results)

Use a tutorial example by Vladislav Korablinov as reference

\*I didn't check if it works with the latest Elasticsearch version

# Main concepts

- Mappings: document structure  
very simple in our case: a single text field
- Analyzers: tokenization, filtering, normalization  
our case: whitespace tokenization, --/+ stemming
- Query-document matching  
BM25 (default)

## Connect

```
es = Elasticsearch([{'host': 'localhost', 'port': 9200, 'timeout': 360, 'maxsize': 25}])
```

## Create index

### Create empty index

```
es.indices.create(index='myindex')
```

### Create index with proper configuration

We are ready to use index setting. Let's define a function which allows us to easily update index settings.

```
def recreate_index():  
    es.indices.delete(index='myindex')  
    es.indices.create(index='myindex', body=settings_final)
```

```
recreate_index()
```

See configuration examples in Vladislav's notebook

## Index documents

At this point we want to add documents to the index. The easiest way to do this is using `parallel_bulk` API. First of all, we have to create a function, which builds an Elastic *action*. *Action* is actually just an index entry, which consist of several meta-fields. We will be focused on 3 of them. `_id` field is literally unique document identifier. `_index` field shows which index the document belongs to. And `_source` field contains document data itself as a JSON object. Let's code it.

```
def create_es_action(index, doc_id, document):
    return {
        '_index': index,
        '_id': doc_id,
        '_source': document
    }
```

Now we have to get some iterable of actions. The most appropriate solution in many cases is creating a generator function. I have my data JSON-represented, so generator will be quite simple:

```
def es_actions_generator():
    for doc_id in range(12):
        with open(f'sample_docs/document_{doc_id}.json', 'r') as inf:
            doc = json.load(inf)
            yield create_es_action('myindex', doc_id, doc)
```

And finally we run indexing.

```
for ok, result in parallel_bulk(es, es_actions_generator(), queue_size=4, thread_count=4, chunk_size=1000):
    if not ok:
        print(result)
```

You can use `elasticsearch.helpers.bulk`

## Search

Here we are, ready to perform search!

We will use `search` API, which takes query as a JSON object and returns a response as a JSON object too. Let's define a pair of useful functions for visualization of results.

```
def search(query, *args):
    pretty_print_result(es.search(index='myindex', body=query, size=20), args)
    # note that size set to 20 just because default value is 10 and we know that we have 12 docs and 10 < 12 < 20

def pretty_print_result(search_result, fields=[]):
    # fields is a list of fields names which we want to be printed
    res = search_result['hits']
    print(f'Total documents: {res["total"]["value"]}')
    for hit in res['hits']:
        print(f'Doc {hit["_id"]}, score is {hit["_score"]}')
        for field in fields:
            print(f'{field}: {hit["_source"][field]}')

def get_doc_by_id(doc_id):
    return es.get(index='myindex', id=doc_id['_source'])
```

See query variants in Vladislav's notebook

# Task (for now)

- Index WikiR **en1k** collection, estimate document indexing time
  - Without stemming
  - With stemming
- Run *test* queries, get top20 results for each query, estimate query execution time
- Save triples  
<queryID, docID, score>  
for two variants
- \*add one more variant:
  - Lemmatized collection (don't forget to lemmatize queries)
  - Boost phrase matches

# Alternatives

- gensim  
<https://github.com/RaRe-Technologies/gensim/pull/3304>  
not well documented, you have to take care of doc ids
- sklearn.feature\_extraction.text.TfidfVectorizer  
tf.idf only (not BM25)
- [https://github.com/dorianbrown/rank\\_bm25](https://github.com/dorianbrown/rank_bm25)  
presumably slow at query time
- <https://github.com/AmenRa/retriv>  
looks fine, but not matured yet, I guess



# Part 2: Evaluation

- Format your runs in TREC format
- Use [https://github.com/terrierteam/ir\\_measures](https://github.com/terrierteam/ir_measures) (more formats, better documentation)  
or  
[https://github.com/cvangysel/pytrec\\_eval](https://github.com/cvangysel/pytrec_eval) (is in fact behind ir-measures)
- Calculate  $p@10$ ,  $p@20$ , and MAP for your runs and the dataset creators' BM25 run (see test/BM25.res)

# TREC formats

## qrels

158491	0	2102124	1
158491	0	2413096	1
158491	0	785032	1
158491	0	2416831	1
158491	0	1990243	1
5728	0	5728	2
5728	0	957396	1
5728	0	737951	1
5728	0	375146	1
q_id	not_used	doc_id	relevance_label

## Runs

158491	Q0	625257	0	15.660703104969318	BM25
158491	Q0	663828	1	15.576630390508356	BM25
158491	Q0	607552	2	15.42499982440102	BM25
158491	Q0	93661	3	14.900135903438647	BM25
158491	Q0	1902136	4	14.900135903438647	BM25
158491	Q0	1490799	5	14.852102590235583	BM25
158491	Q0	1422090	6	14.824568369009627	BM25
158491	Q0	1880296	7	14.710114506753003	BM25
158491	Q0	2261272	8	14.710114506753003	BM25
158491	Q0	13801	9	14.515017321771746	BM25
158491	Q0	621578	10	14.515017321771746	BM25
158491	Q0	635537	11	14.074421958858867	BM25

q_id	n_u	doc_id	rank	score	run_name
------	-----	--------	------	-------	----------